Reinforcement Learning: An Introduction
Attempted Solutions
Chapter 1

Scott Brownlie & Rafael Rui

1 Exercise 1.1: Self-Play

Suppose, instead of playing against a random opponent, the reinforcement learning algorithm described above played against itself, with both sides learning. What do you think would happen in this case? Would it learn a different policy for selecting moves?

dwd